

Your grade: 100%

Your latest: 100% • Your highest: 100% • To pass you need at least 80%. We keep your highest score.

[Next item →](#)

1. Which techniques does quantized low-rank adaptation (QLoRA) use for fine-tuning large language models (LLMs)?

1 / 1 point

- ☐ LoRA adaptation
- ☒ 4-bit quantization
- ☐ Few-shot inference
- ☐ Zero-shot inference

✓ Correct

4-bit quantization maintains high precision in the models. It reduces memory usage, making the model highly efficient for fine-tuning LLMs.

2. Which of the following techniques helps reduce the number of trainable parameters by adding low-rank matrices?

1 / 1 point

- ☐ Additive fine-tuning
- ☐ Soft prompts
- ☐ Full fine-tuning
- ☒ LoRA

✓ Correct

Reparameterization-based methods, such as LoRA or low-rank adaptation, use reparametrizing network weights using low-rank transformations. This reduces the number of trainable parameters while still working with high-dimensional matrices like the pre-trained parameters of the network.

3. LoRA helps fine-tune a pre-trained language model. If you introduce low-rank matrices to the model, how does it affect its parameter efficiency?

1 / 1 point

- ☒ Introducing low-rank matrices to the existing weights helps add a small number of trainable parameters.
- ☐ Low-rank matrices keep track of the original number of parameters and use them for an efficient structure.
- ☐ Low-rank matrices replace the original weight matrices with low-rank approximation to reduce the number of trainable parameters.
- ☐ Low-rank matrices increase the number of trainable parameters by adding high-weight matrices.

✓ Correct

Introducing low-rank matrices to the existing weights makes the fine-tuning process parameter-efficient and increases the model size.

4. Which of the following code snippets copies the original linear model and creates a LoRALayer object?

1 / 1 point

☐

```
class LoRALayer(torch.nn.Module):
    def __init__(self, in_dim, out_dim, rank, alpha):
        super().__init__()
        std_dev = 1 / torch.sqrt(torch.tensor(rank).float())
        self.A = torch.nn.Parameter(torch.randn(in_dim, rank) * std_dev)
        self.B = torch.nn.Parameter(torch.zeros(rank, out_dim))
        self.alpha = alpha

    def forward(self, x):
        x = self.alpha * (x @ self.A @ self.B)
        return x
```

☐

```
from urllib.request import urlopen
import io

model_lora = TextClassifier(num_classes = 4, freeze = False)
```

```
model_lora.to(device)
path = 'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/ \
      uGC04Pom651hQs1XrZ0NsQ/my-model-freeze-false.pth'
urlopened = urlopen(path)
stream = io.BytesIO(urlopened.read())
state_dict = torch.load(stream, map_location = device)
model_lora.load_state_dict(state_dict)
```

```
model_lora.fc2 = nn.Linear(in_features = 128, out_features = 2, bias = True).to(device)
model_lora.fc1 = LinearWithLoRA(model_lora.fc1, rank = 2, alpha = 0.1).to(device)
```

```
class LinearWithLoRA(torch.nn.Module):
    def __init__(self, linear, rank, alpha):
        super().__init__()
        self.linear = linear.to(device)
        self.lora = LoRALayer(
            linear.in_features, linear.out_features, rank, alpha
        ).to(device)


    def forward(self, x):
        return self.linear(x) + self.lora(x)
```

✓ Correct

LinearWithLoRA applies the original linear model and LoRA model to the input X and adds the output together.

5. What would be the equation for the resultant parameter when the LoRA layer's output computes as a function of $h(x) = W_0 \times x$, where x = input from the last layer?

1 / 1 point



- ☐ $d \times x$
- ☐ $k \times x$
- ☐ $d \times W_0$
- ☒ $d \times k$

✓ Correct

The original neural network layer has a weight matrix W_0 with dimensions d by k , where d is the input size and k is the output size.